

HOW NETFLIX AUTOSCALES CI

Rahul Somasunderam

WHAT DOES CI LOOK LIKE AT NETFLIX

JENKINS @ NETFLIX

- 35 Jenkins controllers
- ~45k job definitions
- ~600k builds per week
- 650-1500 agents
- 1-100 executors per agent

THE SPINNAKER VIEW

- 1 Application
- 35 stacks (Controller Clusters)
- 180 Agent Clusters
- 1+ ASG per cluster
- All workloads on AWS

CLUSTERS AND ASGS

CLUSTERS AND ASGS

- AWS has Auto Scaling Groups

CLUSTERS AND ASGS

- AWS has Auto Scaling Groups
- Spinnaker calls them Server Groups

CLUSTERS AND ASGS

- AWS has Auto Scaling Groups
- Spinnaker calls them Server Groups
- <Application>-<Stack>-<Detail>-v<Version>

CLUSTERS AND ASGS

- AWS has Auto Scaling Groups
- Spinnaker calls them Server Groups
- <Application>-<Stack>-<Detail>-v<Version>
- jenkins-unstable-agent-highlander-v123

HOW TO PLAN FOR CI INFRASTRUCTURE

INFINITE RESOURCES

- Provision capacity based on known maximum load
- Multiply by a safety factor for good measure
- Monitor and change the capacity as load increases

INFINITE PATIENCE

- Plan capacity based on median load
- Builds will sit in queue for long times

INSTANT RESOURCES

- You will get resources as soon as you request for them
- Works well with Containerizable builds
- Not all builds can be containerized
- Does not scale well with large numbers of short-lived builds

AUTOSCALING

- Set up minimum and maximum capacity
- Scale based on some metric

WHAT METRIC TO USE

SYSTEM METRICS

SYSTEM METRICS

CPU/Memory/Disk IO/Network throughput

- Natively supported by cloud providers and most metrics solutions

SYSTEM METRICS

CPU/Memory/Disk IO/Network throughput

- Natively supported by cloud providers and most metrics solutions

Scaling Policies are supported by cloud providers

SYSTEM METRICS

Not very useful for CI

QUEUE DEPTH

QUEUE DEPTH

Queue Depth seems adequately proportional.

QUEUE DEPTH

Queue Depth seems adequately proportional.
However, it is a trailing metric.

AGENT UTILIZATION

AGENT UTILIZATION

For each agent, find [idle, busy, offline] executors.

AGENT UTILIZATION

For each agent, find [idle, busy, offline] executors.

Sum these up by ASG.

AGENT UTILIZATION

For each agent, find [idle, busy, offline] executors.

Sum these up by ASG.

Compute utilization as $\frac{busy + offline}{busy + offline + e}$

MEASURING AGENT UTILIZATION

AN AGENT'S ASG

When launching agents, use labels to specify the placement of the agent.

 Agent **nflux-agent-unstable-i-0522989245ff3659d** (Connect: `ssh -t i-0522989245ff3659d`)

Mark this node temporarily offline

Agent is connected.

Labels

```
asg:jenkins-unstable-bionic-v189 aws:test:us-east-1:jenkins-unstable-bionic-v189 bionic buildgroup:bionic carson.version:0.767.0 carson:true cloud:aws cluster:jenkins-unstable-bionic detail:bionic ec2.availZone:us-east-1e ec2.instanceType:m5d.xlarge ec2.region:us-east-1 env:test executors:4 iamRole:jenkinsInstanceProfile java.jvm:zulu8 java.runtime:1.8.0_292-b10 nf.account:test nf.app:jenkins nflux.agent.build:569 os.arch:amd64 os.codename:bionic os.distribution:ubuntu os.name:linux os.release:18.04 stack:unstable us-east-1
```

CAPTURING METRICS

We wrote a custom plugin that plays well with Atlas.
You could write one for whatever your metrics
capturing service is.

AUTOSCALING

HOW TO AUTOSCALE

AWS offers 2 ways to scale

- Target Tracking
- Step Scaling

WHEN TO SCALE UP

Edit scaling policy X

Conditions

Whenever of [Search all metrics](#) ?

is None

for at least consecutive period(s) of

HOW TO SCALE UP

Actions

Add	▼	20	percent of group ▼	when jenkins.executorsUtilization is between 0.65 and	0.8
Add	40	percent of group	when jenkins.executorsUtilization is greater than or equal to 0.8		
Add step					

Documentation

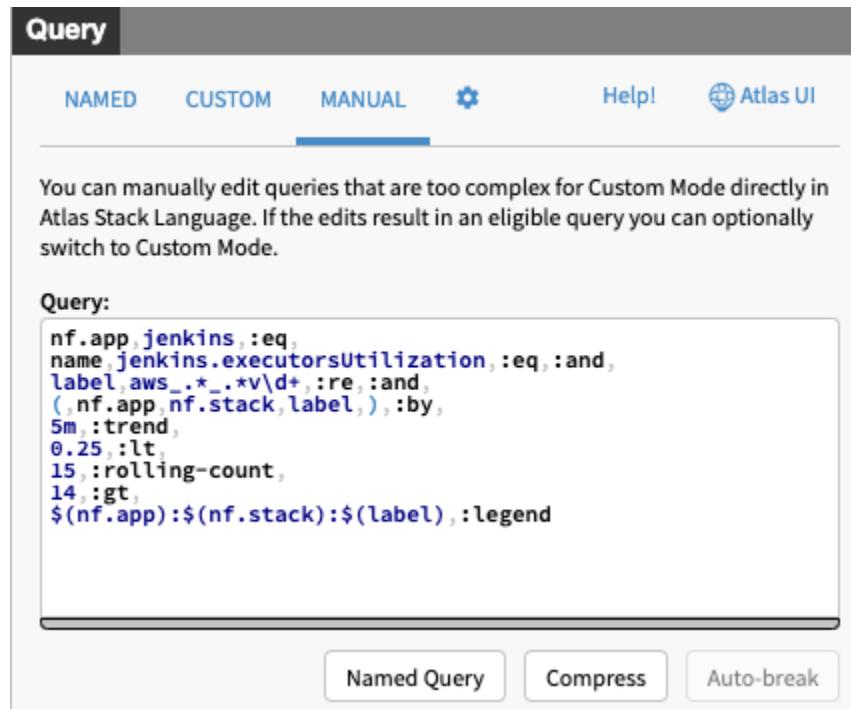
Additional Settings

Policy Name jenkins-buildstest-bionic_classic-v030-NFLX/EPIC-jenkins.executorsUtilization-GreaterThanThreshold-0.65-1-60-1620084562030

Adjustment Step Add instances in increments of at least instance(s)

Warmup Instances need seconds to warm up after each step

WHEN TO SCALE DOWN



The screenshot shows the MongoDB Atlas Query interface. The top navigation bar has tabs for 'NAMED', 'CUSTOM', 'MANUAL' (which is underlined), 'Help!', and 'Atlas UI'. Below the navigation is a descriptive text: 'You can manually edit queries that are too complex for Custom Mode directly in Atlas Stack Language. If the edits result in an eligible query you can optionally switch to Custom Mode.' A large text area labeled 'Query:' contains the following MongoDB Stack Language query:

```
nf.app,jenkins :eq,  
name,jenkins.executorsUtilization :eq,:and,  
label,aws_*.*v\d*:re,:and,  
(,nf.app,nf.stack,label):by  
5m:trend  
0.25:lt  
15:rolling-count  
14:gt  
$(nf.app):$(nf.stack):$(label):legend
```

At the bottom of the query editor are three buttons: 'Named Query', 'Compress', and 'Auto-break'.

HOW TO SCALE DOWN

RECAP

WHAT WE LEARNT

WHAT WE LEARNT

- This improved support experience

WHAT WE LEARNT

- This improved support experience
- This improved the experience for spiky workloads

THANK YOU!

jobs.netflix.com